LARGE-SCALE GRAPH COMPRESSION VIA ASYMMETRIC NUMERAL SYSTEMS

SUPERVISOR:

Sebastiano Vigna

CO-SUPERVISOR:

Paolo Boldi

CANDIDATE:

Francesco Tomaselli

Matr. 969853

Academic year 2021/2022

Abstract

This thesis focuses on the application of asymmetric numeral systems to large-scale graph compression.

Asymmetric numeral systems are a family of entropy encoders that obtains compression quality comparable to arithmetic coding, thus optimal for a given source of symbols, while having a decoding speed similar to Huffman coding. They were presented by Jarek Duda and are heavily used in several different scenarios, such as Linux kernels, Facebook Zstandard, JPEG XL and many others.

This presents an application of such encoders on large-scale web and social graphs. The final proposed methodology is the result of three iterations and combines asymmetric numeral systems, instantaneous codes and patched frame of reference encoding.

The experimental results show how this methodology saves as much as 76 percent of space with respect to quasi-succinct representations and is capable of storing graphs in as low as 3.5 bits per link.

Contents

Chapter 1

Introduction

Graphs are one of the most versatile data structures in computer science and mathematics, finding an extreme variety of applications from path-finding, to ranking in search engines, community detection, and so on. One of the key requirements of graph analysis is the possibility to efficiently visit them, which is an easy task when they are small, but not so much when billions of nodes and hundreds of billions of arcs are involved. A web graph is a structure where nodes are made out of web pages, and there exist a link between nodes x and y if x contains a hyperlink to page y. On the other hand, social graphs are structures where nodes are connected via a social relation, for instance Wikipedia pages, friends and followers on a social network, and so on. These types of graphs certainly store a large amount of information, thus we need a way to visit them efficiently.

Why the classical representation do not work The classical ways of storing graphs are through adjacency matrixes or lists: in the first case a $n \times n$ square matrix is considered and a boolean flag on cell i, j indicates whether there exists an arc between node i and j, adjacency lists represent a graph through a list of successors for each node. The classic representations are perfectly fine for small graphs, but occupy a huge amount of memory for large-scale ones. In fact, one of the main problems associated with web and social graphs is their dimensions, we are talking about structures with possibly billions of nodes and hundreds of

billions of arcs, thus a classical representation can not be loaded in memory. A possible workaround to visit a large-scale graph is to load smaller parts of the structure from disk when needed, but this is slow and unpractical.

Data compression   One solution to visit large-scale graphs is data compression. The goal is to store structures in fewer bits than a classical representation, and this often translates to finding ways to compress integers. Data compression makes it possible to load larger structures in memory but also benefits the scenario where we are forced to partially load them from disk, as each partial load considers a greater percentage of the total structure.

There exist many techniques for data compression, and in the end, they all involve integer compression. Some classical approaches are found in the application of instantaneous codes to encode a given integer source. Those have been widely used to store all sorts of data and are straightforward to apply. The idea is to encode an integer n in a codeword, where the length of it determines its probability, hence one needs to assign shorter codewords to frequent integers.

Outside of the realm of instantaneous codes, we find techniques to efficiently code blocks of integers, such as patched frame of reference, and structures that use a number of bits close to the succinct bound such as Elias-Fano lists.

Other techniques do not rely on an intrinsic distribution of codewords and are tailored to specific sources, such as asymmetric numeral systems, that are a family of entropy encoders. They use optimal space for a given source, paying the price of the need to store additional information. This type of encoders are the main focus of this thesis, indeed we want to find out if they can be fruitfully applied to graph compression.

Graph compression has been already studied in the past, a good example is the Webgraph framework, which exploits the structural properties of web and social graphs to store them in a few bits per link. The approach relies on the empirical similarity of successors of close nodes, using reference compression to represent a node's outgoing links: the idea is to fix a reference node and represent its successors with a bitmask. The residuals, which are nodes not included in the reference, are

written exploiting consecutivity, which creates several consecutive intervals in a node successors list.

Another starting point for this study is the compression of inverted indexes. An inverted index is a data structure that builds the backbone of search engines, the idea is to assign to each word the document ids, plus other additional information, where it appears. As the reader can imagine this structure is similar to an adjacency list, as we are storing a list of integers for each term. Many techniques can be applied to store this type of index, but one work, in particular, exploits asymmetric numeral systems for the task. The authors suggest some key disadvantages to keep into consideration when applying ANS to large-scale datasets.

The proposed methodology The proposed methodology is a compression scheme for graphs, where asymmetric numeral systems are applied to the gap encoding of the nodes' successors. We start with a naïve application to then refine it. One of the characteristics of entropy encoders is that they can be tailored to a statistical source, hence they offer optimal compression ratios, the tradeoff is that we need to store additional information to rebuild these ad hoc compressors. Indeed, ANS in particular uses symbols maps to store observed probabilities, and they can become quite large if a source exhibits several different symbols.

The naïve approach fails exactly for this reason, as we are building a specialized model for each node in a graph. The technique generates optimal compression while wasting an enormous amount of space writing all the additional information required by each encoder on disk. To solve this problem, various approaches are employed, but the idea is to first limit the size of the symbols maps of the encoders. One way to do so is through symbols escaping, which consists in removing some symbols from the maps and encoding them with another compression method. The other point left to solve is the excessive amount of specialized models, as we are dealing with a different encoder for each node. The idea that comes to mind is clustering, which consists in finding some similarity relations between models, reducing a group of similar ones to a single encoder. For this purpose, a Gray code ordering is employed on the models' symbols, to then partition them heuristically.

Finally, the idea of escaping is applied to the clustered models, where we find a suitable set of symbols to remove from the frequency maps by minimizing a space estimation function.

The final result is a methodology that combines asymmetric numeral systems, patched frame of reference, and instantaneous codes to store large-scale graphs. The results are promising and the compression schema is capable of storing web graphs in as low as 3.5 bits per link, saving 76 percent of the space required by quasi-succinct data structures. The access speed to the graph is also comparable to the existing methods, keeping in mind that entropy encoders are on the slower side of compression methods.

Experimental results and conclusions   All the implementation is written in Java and publicly available. The project implements Webgraph interfaces making it easier to store and load graphs from the public datasets. Compression results are expressed in bits per link, and are always in between of the Webgraph methodology and the quasi-succinct representation. Speed is tested by randomly accessing nodes successors, and the obtained speed per link is comparable to the Webgraph compression scheme.

Dissertation organization   The dissertation is organized as follows: Chapter 2 on the facing page gives an overview of different compression techniques, starting from instantaneous codes for integers, ending with asymmetric numeral systems, patched frame of reference, and Elias-Fano monotone lists. Chapter 3 on page 13 introduces the Webgraph compression scheme and a relevant study about inverted index compression. In Chapter 4 on page 21 the proposed methodology is discussed, explaining the challenges, the refinements, and the final approach to graph compression. Chapter 5 on page 42 presents all the experimental results. We start by defining the graphs used in the experiments, which come from the Webgraph archive and are a mix of web and social ones, to then proceed with the compression results on all those graphs and the speed of access tests.

Chapter 2

Compression techniques

We now introduce various compression techniques. The first section talks about instantaneous codes, which are a way to encode an integer source. Later, other techniques are presented, namely asymmetric numeral systems, patched frame of reference, and Elias-Fano monotone lists.

2.1 Instantaneous codes

Instantaneous codes have been widely used to store integer sources. The idea is to assign a codeword to each integer, the length of the word determines a probability distribution, and thus each code is better suited to a given integer source. Other applications can be found in inverted indexes storing, a structure where for each word we store the documents where it appears. A common use of codes is to pre-process the document list by gap encoding and then store the gaps with some code on disk.

What is a code   A code is a set $C \subseteq 2^*$ of binary words. We define an ordering by prefix of the words in 2^* as the following:

$$x \preceq y \iff \exists z \mid y = xz$$

Two words are incomparable if neither is a prefix of the other.

A code is said to be instantaneous if each pair of distinct words in the code are incomparable. This property assures that every given word w obtained by concatenation of words in the code is uniquely decodable. An example of instantaneous code is $\{0, 1\}$; $\{0, 01, 1\}$ is not. Considering $w = 01001$, the first code decodes it successfully, the second creates ambiguity between the words 01 and 0.

A code is complete if every word $w \in 2^*$ is comparable with some word of the code. If this property holds, no words can be added to the code without violating instantaneousness.

Kraft-McMillan inequality   Let $C \subseteq 2^*$ be a code. If C is instantaneous, then:

$$\sum_{w \in C} 2^{-|w|} \leq 1$$

A code C is complete if and only if the equality holds. Also, given the, eventually infinite, sequence $t_0, t_1, \ldots, t_{n-1}, \ldots$, if this equation holds

$$\sum_{i=0}^{n} 2^{-t_i} \leq 1$$

there exists an instantaneous code made up of words $w_0, \ldots, w_n$, such that $|w_i| = t_i$.

2.1.1   Codes examples

Unary code   This is one of the simplest codes, indeed an integer n is written as a sequence of n zeros followed by a single one. This code is not optimal to use as is but is an integral part of more performing ones. The length of the codeword associated with n is trivially $n + 1$.

Note that zeros and ones could be swapped leading to the same result, this would create consistency between integer and lexicographic ordering of the codewords.

Elias γ code   A more sophisticated encoding ig γ, presented by Peter Elias [10]. This time an integer n is encoded in a *length, payload* fashion. The latter is

computed with the binary representation of n minus the leading one, for example:

$$123 = 1111011 \rightarrow 111011$$

γ then writes the length of the payload in unary, followed by it, thus:

$$123 \rightarrow 0000001|111011$$

It follows that for a generic integer n, the length of the code is $2\lfloor \log n \rfloor + 1$.

Elias δ code   The idea here is the same as γ, but this time we replace the unary code to write the length of the payload with γ itself [10]. Hence the length of the codeword for n is $2\log(\log(n+1)+1)+1+\log(n+1)$.

Note that we could go on in this way and define a code that uses δ to write the payload length, but there are no significant advantages.

Golomb code   This code fixes a base b and writes n as the quotient of n/b in unary, followed by $\log b$ bits for the remainder [11]. The length of the codewords are $\lfloor \frac{n}{b} \rfloor + 1 + \log b$.

A note about this code is that, if we fix the base at

$$b = \left\lceil -\frac{\log(2-p)}{\log(1-p)} \right\rceil$$

we create an optimal code for a geometric distribution of parameter p.

Varyable byte codes   This type of representation encodes an integer in some number of bytes [18]. The idea is to consider a single byte as a *continuation bit* plus a *payload*, containing part of an integer. To encode an integer we write it as blocks of 7 bits, where all but the last one have concatenation bit set to 0.

To decode, we iterate over the blocks until a continuation bit of 1 is reached, all the 7-bit parts are then merged through bit manipulation. The size of the blocks creates a tradeoff between compression and decoding speed, some popular choices can be 32, 16, 8 or 4 bits blocks.

n	γ	δ	Golomb	VByte
0	1	1	100	1000
1	010	0100	101	1001
2	011	0101	110	1010
3	00100	01100	0100	1011
4	00101	01101	0101	1100
5	00110	01110	0110	1101
6	00111	01111	00100	1110
7	0001000	00100000	00101	1111
8	0001001	00100001	00110	00011000
9	0001010	00100010	000100	00011001
10	0001011	00100011	000101	00011010

Table 2.1: Codewords examples. Golomb code uses $b = 3$, and VByte has a block size set to 4.

2.2   Other techniques

We now present other compression techniques that do not fall into the realm of instantaneous codes. The first one is asymmetric numeral systems, the second is a methodology to store a list of integers exploiting binary magnitudes of integers, and the third is a quasi-succinct data structure for lists. Note that many other techniques exist, for the sake of brevity we present here only those that are relevant to this study.

2.2.1   Asymmetric numeral systems

Asymmetric numeral systems are a family of entropy encoders presented by Jarek Duda that achieves a compression ratio comparable to arithmetic coding while having a decoding speed similar to Huffman coding [6][7][8]. They are used in a wide range of applications, for instance, Facebook Zstandard compressor, JPEG XL, Linux kernel and so on.

Standard numeral systems   In standard numeral systems, symbols are treated as if they carry the same amount of information. Fixing the binary case, we have

two symbols, 0 and 1, each with a probability of $1/2$, thus the optimal space required to store each becomes $\log(1/2)$. If we think about it, this is exactly what happens in the classical binary encoding, given x as the current sequence, it becomes $x' = 2x$ or $x' = 2x + 1$, if we add a zero or a one, thus each symbol takes exactly $\log(1/2)$ bits. The idea of Duda is to extend this concept to the asymmetric case, creating a similar scenario where we have many symbols each with a different probability.

Asymmetric case What the binary example shows us is that a single integer holds an encoded sequence, this is also the case for the asymmetric case. The idea is to encode a generic sequence of symbols in a single integer x, which will be called *state*. More specifically, fixed a symbol s that appears with probability p_s, this method obtain a new state $x' \approx x/p_s$. The result is an encoding that uses a number of bits close to the Shannon entropy of the source of the symbols and is thus optimal.

Let S be a symbol source, consider now a generic symbol s and its frequency F_s. We scale all the counts to a fixed range 2^d and obtain new frequencies f_s for each symbol, thus the probability for s is approximated by $f_s/2^d$.

We define the cumulative c_s for each symbol as the quantity:

$$c_s = \sum_{t < s} f_t, \quad f_t > f_s \rightarrow t < s$$

this permits to define the primitive $sym(n) = \max_{c_s \leq n} S$, with $n \in [1, 2^d]$, which is required both to encode and decode a symbol.

Encoding and decoding The state, initially set to zero, is updated at each encoding and decoding, the primitives are defined as:

$$encode(state, s) = \lfloor state/f_s \rfloor * M + c_s + state \bmod f_s$$
$$decode(state) = \langle \lfloor (state - r)/M \rfloor * f_s - c_s + r, s \rangle$$

where $M = 2^d$, $r = 1 + (state - 1) \bmod M$, and the decoded $s = sym(r)$.

For the above definition, the encoded sequence is decoded in reverse order. In practice, once the state overflows, it is normalized by storing the current state and resetting it to zero. This worsens the compression ratio but permits handling integers with a fixed length.

Example   Lets assume to have a list of symbols:

$$1, 1, 1, 2, 2, 1, 2, 3, 1, 1, 3, 2, 1, 2, 3, 1$$

we can compute the probability of each symbol and scale it to a fixed range, assume it to be 1024, the cumulative can also be computed:

Sym	F_s	f_s	c_s
1	8/16	512	0
2	5/16	320	512
3	3/16	192	832

We can now start to encode the symbols:

$$\text{encode}(0, 1) = (0/512) * 1024 + 1 + 0 = 1 \qquad \text{(1bit/1)}$$

$$\text{encode}(1, 1) = (1/512) * 1024 + 1 + 1 = 2 \qquad \text{(2bit/2)}$$

$$\text{encode}(2, 1) = (2/512) * 1024 + 1 + 2 = 3 \qquad \text{(3bit/3)}$$

$$\text{encode}(3, 2) = (3/320) * 1024 + 513 + 3 = 516 \qquad \text{(11bit/4)}$$

$$\text{encode}(516, 2) = (516/320) * 1024 + 513 + 196 = 1733 \quad \text{(12bit/5)}$$

$$\text{encode}(1733, 1) = (1733/512) * 1024 + 1 + 197 = 3270 \qquad \text{(13bit/6)}$$

$$\text{encode}(3270, 2) = (3270/320) * 1024 + 513 + 70 = 10823 \quad \text{(15bit/7)}$$

$$\cdots$$

The result is an integer of 30 bits storing 16 elements, overall less than 2 bits per element are used.

2.2.2 Patched frame of reference

This mechanism operates in the context of storing elements in a block of size B. Let b_i be the i-th element of the block, one could use a fixed representation set at $\max_{i=0,\dots,B} \lceil \log b_i \rceil$ bits to store every element. This is perfectly fine but it might be a waste of space in cases where a few outliers are present. For instance, given the block of 8 elements:

$$12, 23, 12, 1, 2, 176432, 3, 2$$

the method would use $\lceil \log 176432 \rceil = 18$ bits for each element, when we can clearly see that $\lceil \log 23 \rceil = 5$ bits would be enough to store all elements apart from the maximum.

We can build on top of this idea and compute a binary magnitude l that covers most but not all elements in a block b. The residuals are written in another way at the end of the block. To determine l a search over likely values is performed and the one leading to better compression is selected. Many variants exist that encode exceptions in various ways, preferring speed, or compression ratio [13][21].

2.2.3 Elias-Fano monotone lists

This method was proposed by Peter Elias and it consists of a quasi-succinct representation of monotone sequences [9]. Quasi-succinct structures use a space that is close to the information-theoretical lower bound and provides on average constant-time access to the stored data. In a more recent paper, this structure is used to store inverted indexes, obtaining good results both in compression and speed [20].

High/low bits representation Assume to have a monotonically increasing sequence of n values that are limited by an upper bound u:

$$0 \leq x_0 \leq x_1 \leq \cdots \leq x_{n-1} \leq u$$

Each element is split into high and low bits as follows:

- the lower $l = \max\{0, \log(u/n)\}$ bits are written in a *lower-bits array* sequentially;

- the upper bits are stored in the *upper-bits array* written by gap in unary code.

The upper/lower bits techniques uses at most $2 + \lceil \log(n/u) \rceil$ bits per element. The information-theoretical lower bound for a monotone list of n element in a universe of u elements is

$$\left\lceil \log \binom{u+n}{n} \right\rceil \approx n \log \left(\frac{u+n}{n} \right)$$

Elias proved that this representation is really close to the succinct bound, hence we can call it quasi-succinct.

Primitives   To get the i-th element from the list, we perform i unary code reads to obtain the upper part of the element, combining them with the lower bits found on position i in the lower array. One can speed up the search by fixing a quantum q and storing in a table the position of the q-th zero or one. This permits skipping some parts of the list and completing the remaining part of the search sequentially.

Chapter 3

Graphs and inverted indexes

This chapter deals with graph and inverted index compression. The two tasks are closely related: in a graph one must encode the successors' ids for each node, while in an inverted index each word is associated with the id of the documents where it appears.

The first section introduces WebGraph, a compression method for large-scale graphs coming from the web. The second one talks about relevant applications of asymmetric numeral systems to inverted indexes compression.

3.1   WebGraph

We now introduce an established graph compression methodology for web graphs that makes heavy use of their properties to store them in a highly compressed way [2][3]. This is the technique on top of which the proposed methodology is built and is referred to from now on as the BV graph compression scheme.

What is a web graph   A web graph is a graph where there exists a node for each URL, and an arc between nodes x and y if a hyperlink in x leads to y. Such graphs can be really useful for a variety of reasons, for instance, one could improve search engine performances, detect cyber communities, improve crawlers, and so on.

The problem that rises when dealing with such structures is the huge cardinal-

ity. We are talking about billions of nodes and hundreds of billions of arcs. Such cardinalities make it impossible to hold the entire graph in memory, hence some sophisticated techniques are required to visit it. Although streaming could be an option, that means loading a particular part of the graph from disk only when necessary, data compression can be employed, decoding successors from a compressed representation that lives in memory instead of disk. Note that compression can also benefit streaming, as each partial load can gather a bigger part of the on-disk structure.

Web graph properties   Assume for now that web graphs are ordered in a lexicographic way by node URL. Some empirical properties can be observed when dealing with ths type of structures, that are *locality*, *similarity* and *consecutivity*.

The first one refers to the fact that, on a given web page, most of the hyperlinks are navigational, i.e., they link a page to others on the same website. This property ensures that, by ordering nodes in a lexicographic way, many arcs are between close ones.

The second property means that pages close to each other exhibit plenty of common successors. This is a consequence of locality, as close pages are usually of the same website, thus they share a great number of navigational links. This property is really strong and usually consecutive nodes share almost all successors of nothing at all.

Consecutivity is a property for which numerous successors for a node are consecutive, so the adjacency list for a URL shows various intervals.

Differences between social and web graphs   While web graphs are a great source of information, the compression of social graphs is also relevant. A social graph is a graph where nodes are connected by some sort of social relation, many examples range from social networks to Wikipedia pages.

The compression of those graphs is more complicated as the structural properties introduced for web graphs are less pronounced. Also, the introduced empirical properties are valid if web pages are sorted in lexicographic order, and this ordering is not present in social graphs. Therefore, ordering the nodes accordingly can

have a major impact on compression.

Datasets The Webgraph framework permits the publication of multiple big graphs in the BV representation[1]. They come from various crawling activities and are a mix of social and web graphs of different dimensions [4]. For each crawling result, both the normal and transposed graph are present.

The proposed methodology loads big graphs through this framework, to then re-store them via asymmetric numeral systems.

3.1.1 Compression format

Suppose we have a web or social graph where nodes are assigned labels from 0 to $N-1$ and consider a lexicographic ordering by URL for web graphs and a suitable one chosen for social graphs. Let us consider the list of successors of a node x, $S(x)$. A first naïve representation would be to store them sorted. This is not optimal as no empirical property is exploited.

Computing gaps A first optimization is to compute gaps between them, where the first successor of x is written as $s_1 - x$, and each subsequent one as $s_i - s_{i-1}$. Note that the first one might be negative, thus it is encoded as $2x$ if greater than zero, or $2|x| - 1$ otherwise.

Copy list A second optimization is to exploit similarity, the idea is to write a list of successors of a node as the difference between a preceding list. Fixed a predecessor y, we write the successors of the node x as a *copy list*, that has a 1 or 0 in the i-th position whenever $S(x)$ contains the i-th successor of y. Some nodes might not be represented this way, thus they are written as a list of extra nodes. The choice of the reference node is made inside a window of size W, which creates a tradeoff between retrieval speed and compression quality. Also, the number of hops is considered: to retrieve a node successors one need to decompress the reference list, that might lead to another reference decompression, and so on.

---

[1]`https://law.di.unimi.it/datasets.php`

Residuals intervals   The extra nodes are written exploiting consecutivity. Assuming that they present a list of intervals, each one of them is written as its left extreme plus its length. This creates again a list of residuals that are coded using differences.

3.1.2   Re-ordering web graphs

Until now, we assumed that web graphs are ordered via lexicographic ordering on node URLs, this creates some interesting properties such as similarity and consecutivity but perhaps is not the best ordering to employ. The ordering of the nodes in a crawled graph can have a major impact on the resulting compression ratio, indeed, many studies tried to tackle this problem and find a reordering technique that maximizes compression. The key points are that an ordering should create as much similarity and consecutivity as possible, as those are the two most useful properties for reference compression and interval encoding.

We start by discussing some preliminary studies made on graph ordering, to then introduce the ordering criteria used by the BV compression scheme [1][5].

Gray code and ordering   An n-bit Gray code is an arrangement of all the binary vectors of length n, such that any two successive elements differ only by one bit. An example of Gray code for $n = 3$ is the following:

x	$\overline{x}$
000	000
001	001
010	011
011	010
100	110
101	111
110	101
111	100

Gray codes can be used in this context to order the adjacency matrix rows, this creates slow changes from row to row, possibly increasing similarity. Permuting

the entire adjacency matrix is unfeasible and for a web graph is even really sparse, so the idea is to compute gray ordering by iterating through two node successors iterators. This can be implemented with an ad hoc comparator, that is going to be adapted for the proposed methodology. Its modified definition is presented in Algorithm 3 on page 34.

Preliminary studies on ordering   Numerous experiments have been made to assess compression quality of different types of ordering, the tested ones are:

- URL ordering: nodes are ordered in lexicographic order by their url;

- Lexicographic row: given the boolean rows of an adjacency matrix, they are sorted lexicographically;

- Gray: gray ordering is performed on the adjacency matrix rows;

- lhbhGray: keep URLs from the same host adjacency, ordering them via Gray code;

- shbhGray: same as lhbhGray but Gray code are computed only on links to the same host.

The second and third orderings are called intrinsic, as they rely only on the graph structure. The others, which combine intrinsic and host information, are called mixed. Note that on social graphs we are missing host and URL information, thus a clustering algorithm is employed, where two nodes in the same cluster indicate that they have the same imaginary host.

Experiments show how intrinsic orderings work well on transposed graphs while failing to beat the mixed ones on normal graphs. The results on social graphs show some potential but the best results are obtained on graphs coming from the web.

Layered label propagation   The presented preliminary studies show how mixed ordering criteria outperform intrinsic ones, thus information about hosts and URLs is crucial to obtain a good compression ratio. Although the URL can be easily obtained for web graphs, social ones fail to provide such data: clustering

algorithms must be used to build imaginary hosts on social graphs. Some good performing algorithms are based on label propagation, the idea is simple: each node starts with a different label, and, at each round of the algorithm, every node updates its label following some rules. Once no labels are updated the algorithm stops. A classic update rule can be the majority of labels in the neighborhood of a node.

We now present *layered label propagation*, an ordering criteria that improves compression performance on both social and web graphs. The algorithm relies on label propagation clustering, in particular, *Absolute Pott Model*: the update of a node label is done by maximizing a function that takes into consideration both the cardinality of a label in the neighborhood and the density of that particular cluster. The APM algorithm has a hyperparameter γ that regulates the impact of the cluster density factor in the update rule and is called the resolution parameter.

The idea of layered label propagation is to compute different rounds of the APM algorithm with randomly drawn γ, which produces a label assignment for each node. The labels obtained at every round are used to update the final ordering of the nodes, in particular, nodes with the same label are kept close, nodes inside the same cluster are left untouched, and finally, order among clusters is defined by the intrinsic ordering of the labels produced by the algorithm. The idea is to capture structural properties at different resolutions and to combine them.

This ordering procedure produces a graph permutation that exhibits increased similarity and locality. The combination of LLP and the BV compression schema leads to great compression ratios for both web and social graphs.

3.2   Inverted indexes compression

An inverted index is of paramount importance in a text search system. A simple index consists of a set of posting lists, each of which contains a sequence of $d_{t,i}$, that are the ids of the documents that contain the term t.

One of the basic approaches to index compression is to transform each sequence with gap encoding, applying an integer compression technique such as instanta-

neous codes. For this purpose, variable byte codes create a great tradeoff between speed and compression ratio, also Golomb codes can be tailored to a specific geometric distribution that best suit a given index. Patched frame of reference is heavily used on documents id and there exist studies on Elias-Fano lists applied to this domain [17].

One work in particular exploits asymmetric numeral systems on indexes. The idea of the authors is to combine classical methods with entropy coders to potentially obtain better compression ratios [15][16]. They show some interesting ideas and problems to deal with that are at the base of the presented methodology. Another study shows how, paying the price of slower queries, this approach leads to great compression [14].

3.2.1 ANS and packed codes

The mentioned work about inverted index compression with asymmetric numeral systems combines them with a classical compression techniques that is packed codes.

Packed codes A packed code is a technique that exploits any localized consistencies in a sequence. A block size of consecutive values is selected, typically $B = 128$, and represented as B same-width binary values. A selector is considered, stored in the block header, that determines the fixed width. An example is to define for instance 16 selector values,

$$S = 0, 1, 2, 3, 4, 5, 6, 7, 8, 10, 12, 14, 16, 19, 22, 25$$

the block header will contain 4 bits to store l, and that block will represent values up to $2^{S[l]}$. Note that this technique is really similar to patched frame of reference, as the latter encodes integers in the exact same way apart from exceptions.

Packed + ANS This technique combines packed codes and ANS compression. The idea is to keep the selector value and a block size B but this time the selector

refers to a specific ANS model, in the sense of a specific set of symbols frequencies.

This means that to encode a block of B integers, statistics about it are computed and an appropriate model is selected. A trivial selector is the maximum m of the block, which refers to a model with m equiprobable symbols. Note that defining a model with equiprobable symbols permits us to avoid writing symbols on disk to reconstruct the model, as the only thing needed to rebuild it is the value m and a given precision. The bits allocated for the selector determine the maximum number of allowed models, in the original paper, the authors limit themselves to 16 of them. Once a model is selected, all the integers in the block are encoded and the selector and the list of states are written on disk.

Two-dimensional context   An improvement on the previous technique is to use additional information coming from the block as a selector value. The authors refer to this technique as a two-dimensional context. An example can be using both the maximum m and the median k of a block B. This means that an appropriate model would need m symbols where the first k ones share a probability of $\frac{1}{2}$. This improves significantly the one-dimensional context results, and several more models are allowed.

3.2.2   Results and key takeaways

The presented methodology is really interesting and combines two unrelated compression methods to store the inverted index. The compression ratio is good compared to other classical techniques. The access speed is on the slow side, but this is acceptable for entropy encoders.

Some key takeaway from this technique is that the number of ANS models must be somehow limited. The solution of the authors is to rely on a selector value that contains statistics about the blocks to encode. This both limits the number of models and symbols, as the first one is determined by the bits allocated for the selector, and the second by the context used to refer to a single model.

Limiting the number of encoders and their dimensionality is going to be a key factor in the proposed methodology.

Chapter 4

Proposed methodology

We now discuss the proposed methodology for graph compression. The final compression scheme is obtained through various iterations, that start with a basic approach, followed by an optimization, and terminated by the refined final methodology.

The first approach consists of a naïve way of applying ANS to graph compression, presented in Section 4.1 on page 23. The idea is to encode every successors list in the graph with an ad hoc encoder. This means that, given the successors of a node, we compute the frequency of the symbols, where those are the integers appearing in the gap encoding of the list, and we encode them optimally. The result is two separate files, one containing the graph itself, and the other the definition of the models.

Numerous problems arise with this approach and are somehow related to the key takeaways presented in 3.2.2 on the facing page. Recalling the inverted index compression with ANS, the authors tackled two main problems, the first one is the number of different encoders used in the compression, and the other was about limiting the max number of symbols per model. The main drawbacks of the naïve approach are indeed the excessive number of total encoders and the symbols per encoder. If we think about it, we are considering as many compressors as the number of nodes in the graph, that is in the order of the billions, while also allowing all possible gaps to be considered as symbols, that assumes values ranging from

zero to the number of nodes in the graph. This methodology creates an encoding on disk that is heavily dominated by the space taken to represent the models, which is roughly double the size of the encoded successors.

Section 4.2 on page 27 introduces the first step toward the final methodology. The immediate question that arises is how to reduce the space taken by the models. There are numerous ways to do so, but the key point is that although we want to reduce space on the .model file, we also want to preserve the goodness of the successors encoding. Some heuristic approaches have been tested but the one that appeared to be the most proficuous consists of adding escaping to the ANS models, where the idea is to avoid including symbols that are rare in the frequency maps. Starting from node successors, we compute the gap encoding and the statistics about the symbols. We order them and cut off a percentage of the final ones, leaving us with a reduced symbol map. The removed gaps are added as a single special symbol, that will be used to encode them in the state. Once an escaped integer is encoded, it is also written in an escaped list, that will be added to the successors encoding on disk. This approach worsens the graph encoding, as we are adding explicit symbols to be written, but reduces the size of the models considerably, improving the total space. Although better, this approach provides a compression ratio that is similar to the uncompressed representation and is thus not useful, the space taken by the models is surely better than the naïve methodology but still too big.

Section 4.3 on page 31 introduces the final refinement, that focus on limiting the number of models. The first thing that comes to mind is clustering, the goal now is to find a way to merge similar models, so their union can be used to encode more than one node. Initially, we tough about K-means clustering, indeed, there exist some studies about clustering probability distributions that can be exploited in this case [12][19]. The main difficulty of this method is choosing a distance metric and avoiding an extreme situation in which all the models get merged in a single cluster. Unfortunately, this classic clustering algorithm was not easily applicable to this use case, hence it was discarded for a more simple approach. The simplified task now becomes ordering the models in some way and partitioning them, merging

the single partitions. The ordering should put similar models close to each other, and the partitions need to limit the total number of symbols per split. A great way to order models is by Gray-code ordering. The details are discussed in the relative section, but the idea is to sort to the models in a way that subsequent frequency maps differ by at most one symbol. Partitioning is done in a heuristic way, limiting the number of symbols per partition, while keeping the number of clusters low. Finally, the idea of escaping is used on the obtained clustered models, this time in a more sophisticated way. This final refinement reduced the overall space by over 40 percent.

4.1 Initial approach

The initial approach consists of encoding each node with an ad hoc encoder. Although optimal for a single node, this methodology becomes quickly unfeasible for larger graphs. The number of bits needed to write models on disk penalizes the goodness of the optimal encoding since we are writing both too many models and symbols.

Performance varies, but this approach often creates an encoding that exceeds the space taken by quasi-succinct representation, making it not useful. This initial implementation gives an idea about the issues that need to be resolved in later phases.

4.1.1 Computing successors statistics

As stated in Section 2.2.1 on page 8, the idea of an asymmetric numeral system is to encode symbols taking into consideration their probabilities. Hence we need to define what are those symbols in our case and what probabilities we are considering.

The first step to encode the outgoing links of a node is to compute statistics about the gaps. Let *succ* be the array containing the sorted successors of the i-th

node, we compute the array *gaps*, with the following rule:

$$gaps[j] = \begin{cases} succ[j] & \text{if } j = 0 \\ succ[j] - succ[j-1], & \text{otherwise} \end{cases}$$

The symbols coded with ANS are going to be the gaps created by this preprocessing phase, so for each distinct value appearing in *gaps*, we compute its frequency. We then scale these counts to a given range 2^d, obtaining a quantity f_s for each $s \in S$, thus the probability of s can be approximated as $p_s = \frac{f_s}{2^d}$.

The magnitude of the gaps ranges from zero to the number of nodes in the graph, but, considering the consecutivity property, the gap distribution should be concentrated in small values. This is the motivation for the approach, as without this property we would need to hope for a concentration in some values, that might not happen. Note that although a concentration of probability on some gaps would be beneficial for ANS, it is not strictly required that such concentration happens on small gaps. That is not the case for instance with instantaneous codes, which are tailored to tackle monotonically decreasing probabilities of the symbols.

4.1.2   ANS based compression

After computing the frequencies of the gaps of a given node, we encode them with ANS compression.

We define an ANS Model, as the collection of the required data structures to encode a symbols source. This model contains the state storing the encoded symbols, and the frequencies defined in the last section, so an array storing normalized counts and the two mappings, as well as the cumulative.

Preventing overflows   The model state is represented with 64 bits long integer. Recall that the encode primitive takes the older state and starts by multiplying it by $M = 2^d$, and this could cause overflows. To prevent it from happening we perform a 128-bit multiplication between the state and M, checking that the upper 64 bits are all zeros. When inevitably an overflow happens, the current state is

added at the end of a list, resetting it to zero.

Sym primitive The goal of the cumulative if to define this primitive, $sym(n) = \max_{c_s \leq n} S$. While one could store an array of length 2^d and fill each spot with the corresponding cumulative, this task can be solved with an Elias-Fano indexed monotone list using less memory. This implementation supports predecessor queries of the type, given an element x, find the index of the maximum element less than x, that is exactly what you need to find. This is not only really fast but also uses optimal memory. Another approach would be to perform a binary search on the sorted cumulative array, having logarithmic complexity for the *sym* primitive.

Reverse order Given the gaps of the successors of a node, they are encoded in reverse order to avoid inverting the decoded sequence. Once the process is complete, the state sequence and the required structures to rebuild the model are written to disk.

Algorithm 1 Encoding procedure. The fifth line is the one responsible for overflows, when one happens, the state must be added to a list and reset to zero.

1: $state \leftarrow 0$
2: **for** $s \in succ$ **do**
3:    $j \leftarrow state \text{ div } f_s$
4:    $r \leftarrow state \bmod f_s$
5:    $state \leftarrow j \cdot M + c_s + r$
6: **end for**
7: **return** $state$

Algorithm 2 Decoding procedure. In reality the state is not a single integer, thus when the current state reaches zero, the next one must be loaded.

1: $list \leftarrow [\,]$
2: **while** $state > 0$ **do**
3:    $r \leftarrow 1 + (state - 1) \bmod M$
4:    $j \leftarrow (state - r) \text{ div } M$
5:    $state \leftarrow j \cdot f_s - c_s + r$
6:    $list \leftarrow list + sym(r)$
7: **end while**
8: **return** $list$

Model coding on disk Once we encode the successors of a node, we need to write all the required information to decode them on disk. This translates to dumping all the required structures on two separate bit streams.

We start with the list of states produced by the encoding, those are saved on the *.graph* stream. Each node stores, all in γ, its out-degree, the id of the corresponding model, the number of states, followed by the list of states written in binary fixed at 64 bits. The id of the model could be omitted here, as every node has its separate encoder, this is not true in the upcoming sections. The model structures are then written on the *.model* file. We start with the size of the list of symbols, where each one is written and ordered by inverse mapping index. After that, the sorted descending frequencies are written backward by gap. All the integers are written in γ. δ code is also an option, but γ works empirically better in this case. To rebuild the symbols mapping, we just need to read the symbols, associating an incrementing index to each one of them. Frequencies must be filled in inverse order, and the cumulatives are obtained at run time.

Node	Outd.	Model id	N. states	States
...	...	...	...	...
14	5	14	1	1873215
15	2	15	1	4732153
16	130	16	16	1237953, 543843, ...
...	...	...	...	...

Table 4.1: Successors encoding on the *.graph* file

Model id	N. symbols	Inv. sorted syms	Gap sorted freq
...	...	...	...
14	3	1, 5, 3	211, 300, 513
15	1	1	1024
16	6	1, 2, 4, 9, 6, 1234	32, 84, 12, 4, 150, 742
...	...	...	...

Table 4.2: Ans models encoding on *.model* file, notice how the frequencies sums up to 1024, that is the assumed precision.

Performance   As said before, the information is divided into two files, one containing encoded successors and the other information about the model structures. The occupied space is optimal for the first one, but tremendous for the second, as

the required structures contain various symbols. The goal of the next section is to tackle this problem.

Graph	BV	EF	Naïve ANS
cnr-2000	2.897	17.838	17.143
uk-2007-05@100000	2.033	13.7	9.137
uk-2007-05@1000000	1.68	16.479	7.849
en-wiki-2013	13.114	19.500	28.838
in-2004	2.172	18.981	14.799
indochina-2004	1.457	18.469	8.572

Table 4.3: Bits per link obtained by the initial ANS compression method, compared to BV graph and Elias-Fano lists. Interestingly, some graphs are encoded in lower space than EF. Those are all web graphs, hence gaps should be really small without many outliers. Performance on social ones is a different story, as they are stored in almost 1.5 times the space taken by EF.

4.1.3 Approach problems

The main drawback of this methodology is the excessive number of models that we are considering. Not only the space occupied on disk is completely dominated by the model encoding, but also, we can not efficiently load the graph in memory. Indeed, a requirement to efficiently visit a ANS compressed graph is to have the ANS models in memory, avoiding a disk read each time a node's successors are required. With the naïve approach this is not possible.

The immediate goal now becomes finding a way to limit the space taken by the structures both on disk and in memory, and for an initial refinement we think about escaping.

4.2 Escaped ANS

This section focuses on the problem of limiting the space taken by each node model. Defining an encoder the perfect way fails for the mentioned problems about the excessive space of the .model file. The task now becomes finding a way to describe

a model in fewer bits or to make each one of them sub-optimal, by gaining on the disk representation.

One of the first approaches to the problem was to build encoders that are defined approximately, based on the median, third quartile, and maximum of the symbols to encode. Given these statistical measures, all the symbols under the median would share a probability of $\frac{1}{2}$, the ones between the median and third quartile $\frac{1}{4}$, and so on. All the symbol's frequencies and cumulatives can be computed at run-time. This would permit us to reduce the model encoding to just three integers, basically nothing with respect to the successor's file. Although appealing, this method fails to store successors efficiently, as the space taken by those increases massively. Overall, the method is worse than the naïve one. The main issue is that although the assumption that smaller gaps appear more often is true, a model is now considering as many symbols as the maximum observed gap in the successor's list. This means that despite the incredibly compact representation of an encoder, it possibly contains a huge amount of symbols, making each state updates waste many bits. An example can be a gap encoded list of:

$$4189043, 1, 1, 1, 2, 3, 1, 874$$

In this case, we are assuming that the symbols from zero to 4189043 are present in the model, making the encoding of the rarer ones, namely 874 and 4189043, cost too much, as they have a probability of $\frac{0.25}{4188169} \approx 6 * e^{-8}$ each.

Given the gap encoding of a successors list, the big gaps are practically random, so perhaps they can be removed altogether and encoded separately. This is exactly the idea at the base of this section.

4.2.1 Escaping gaps

One of the scenarios in which ANS compression works well is when the overall probability of the alphabet is concentrated on a few symbols, and also the cardinality of the symbols set is not too big. As seen until now, the symbols and their frequencies must be written somewhere to rebuild the decoder, thus an excessive

amount of them makes the really good performance of the ANS encoding useless.

We can build on top of this idea and escape some symbols to encode them in a different way. Given a list of successors gaps as follows:

$$10, 1, 1, 1, 2, 1, 2, 123, 256, 312, 3, 3, 2, 1$$

the frequency map looks like this:

Symbol	Frequency
1	5
2	3
3	2
10	1
123	1
256	1
312	1

We can see some symbols appearing only once, hence we could avoid writing these to the .model file and encode them with an instantaneous code or a different compression method. The idea is to aggregate the rarer symbols into a single one, for instance, discarding only the hapax the map becomes:

Symbol	Frequency
1	5
2	3
3	2
*	4

Frequent symbols are encoded as always, an escaped one is encoded as a *, and written explicitly in an auxiliary list in memory. This list is then going to be written on the .graph file with the states.

4.2.2 Choosing what to escape

In the example above, the symbols appearing once are discarded, in practice, a percentage p is selected. The symbols are sorted in descending order by frequency,

in the case of a tie, the smaller one comes first. Finally, the last $\frac{n}{100} * p$ symbols are escaped. The motivation behind this sorting is that we want to both escape rare symbols and reduce the model encoding space. For the latter reason, smaller symbols are selected in case of a tie, as their codewords in γ are shorter.

Escape encoding   The escaped symbols must be encoded in some explicit way on disk. The chosen approach takes inspiration from PFOR delta. Given v as the list of integers to encode, we find b_{max}, the binary magnitude of the maximum, and we consider b_i, the required bits to store the i-th symbol.

Fixed a b_i, if a number can be written in that number of bits, a zero flag is written, followed by b_i bits representing the integer. If the number overflows the fixed number of bits, an one flag is written followed by b_{max} bits to represent it. For instance, given $b_i = 5$:

$$27 \rightarrow \mathit{flag} = 0, \ \mathit{bits} = 11011$$
$$327 \rightarrow \mathit{flag} = 1, \ \mathit{bits} = 101000111$$

Fixed a b_i, the space taken by the list is computed, and the minimum is selected. The resulting coding on disk is b_i, followed by $b_{max} - b_i$ and the length of the list n, all three written in γ code. The remaining are n integers written as explained above. For example, given the escaped list

$$2103, 43829, 658, 89574, 6431$$

the final encoding on disk looks like this, where the first three integers are written in gamma, and the others in a fixed binary length as just explained:

$$13, 4, 5 | \langle 0, 2103 \rangle, \langle 1, 43829 \rangle, \langle 0, 658 \rangle, \langle 1, 89574 \rangle, \langle 0, 6431 \rangle$$

Performance   This approach cuts the number of symbols to be written, reducing the overall space taken by the graph. The bits required to store only successors increase, as some values might be escaped and written separately. Table 4.4 on the facing page shows compression ratios on example graphs. A positive result of

this approach is that bigger graphs can now be loaded successfully, as the smaller models permit the loading of more node encoders, but the issue is still present on bigger graphs.

Graph	BV	EF	Naïve ANS	Escaped ANS
cnr-2000	2.897	17.838	17.143	16.890
uk-2007-05@100000	2.033	13.7	9.137	8.919
uk-2007-05@1000000	1.68	16.479	7.849	7.476
en-wiki-2013	13.114	19.500	28.838	24.949
in-2004	2.172	18.981	14.799	14.759
indochina-2004	1.457	18.469	8.572	8.305

Table 4.4: Bits per link obtained by the Escaped ANS compression method, compared to previous results. Some graphs are more impacted than others, for instance, we can see a 15 percent improvement on the Wikipedia graph. The difference in web graphs is small. Note that to determine the best p value for escaping we tested values from 5 to 60 in steps of 5.

4.3 Clustered ANS

After reducing the number of symbols per model, the goal is to avoid having a different model for each node. The motivation is that although escaping the models permits to load bigger graphs and saves overall space on disk, the occupied model space is still big, leading to overall poor compression performance.

The question now becomes, how can we merge models and what criterion should one choose to do so. The first task is quite simple, but the latter requires a bit more work. We are now putting aside the escaping procedure presented in the previous section and consider only the full models. The escape idea is reintroduced in a more sophisticated way after the models' reduction.

Merging two models   The first step to clustering is defining how we can merge two models. Recall that the ad hoc encoder symbols are observed from a node successor, later, counts get normalized to a given range 2^d to estimate probability.

The union of two or more models is computed by merging the frequency maps

before normalization, for instance, given two models M_1 and M_2, the frequency map of the union U is computed as follows:

Symbol	M_1	M_2	U
1	10	6	16
2	7	8	15
5	-	1	1
14	-	1	1
38	1	-	1

that corresponds to treating the gaps coming from node 1 and 2 as a single long list. The union map is then normalized as always.

K-means clustering for distributions When thinking about clustering, a classic algorithm that comes to mind is K-means, which requires a distance metric and operates iteratively. The idea is to assign each point to the closest centroid, recomputing them accordingly at each iteration. This method could work in our case, as the data points are the ANS models, and a suitable distance metric can be the Kullback-Leibler divergence between probability distributions, computed by averaging its quantity in both directions. The centroid can then be recomputed by summing the observed symbol frequencies for each model and normalizing them, just as discussed in the above paragraph.

Unfortunately, this method fails to perform well for this problem. The main drawback is that the process is not reliable, leading to an unfortunate situation in which more and more points get added to a single big cluster. In the end, a simpler heuristic approach is selected, removing the need to perform a classical clustering algorithm.

4.3.1   Ordering and partitioning

After defining how two models can be merged, the problem of finding suitable subsets of models to cluster remains open, as the application of K-means clustering fails to perform well for the given task.

Computing all possible subsets would be unfeasible, thus a heuristic approach is employed. The idea is, starting from all the models, to order them and find contiguous partitions to merge. This reduces the problem of choosing a suitable ordering for the models and partition criteria.

The ordering criterion should closely place similar models, while the partition logic must find a tradeoff between the number of symbols per model and the number of models themselves. The first problem is solved with Gray-code ordering, a heuristic split procedure is developed to tackle partitions.

Gray-code ordering   Let k_i be a sorted array of symbols in the i-th model, if we fix an upper bound n, that array can be seen as a binary vector on the possible symbols. For example, given a model with some symbols, $3, 1, 8, 5$, and a limit of 10, we sort the array and imagine a bitmask like this:

$$[3, 1, 8, 5] \longrightarrow [1, 3, 5, 8] \longrightarrow 01010100100$$

If the goal is to minimize the difference in symbols between adjacent models, a Gray ordering is what we need. As we saw in 3.1.2 on page 16, Gray codes find another application in web graphs, in particular, they can be permuted via Gray ordering on the adjacency matrix rows, the idea is to have slow changes in successors when going to the next row. The intuition here is the same, we want slow changes in the model symbols when we go from one model to the following. The mentioned study also proposes an interesting comparator that makes it possible to order two adjacency lists by Gray code, without having a complete instance of the row. We can use the same comparator to order the models by symbols, its definition can be seen on Algorithm 3 on the following page.

Partitioning the sorted models   After sorting the models accordingly, one must find a way to determine how to merge them. The idea of the approach is to specify an upper limit to the number of partitions and let the algorithm do its work, ideally, we want models with a limited and somehow equal number of symbols. Starting from an initial max symbols quantity, which is set to the average

---

Algorithm 3 Gray code comparator for two ANS Models. i and j denote two iterators, returning the model keys in sorted order, where ∞ means that the end is reached. The expression $[a < b]$ has value 1 if a is smaller than b, zero otherwise. The return value of the comparator is less, equal or greater than zero if the j-th model precedes, is equal or follows i-th model.

1: $p \leftarrow \textit{false}$
2: **while** $\textit{true}$ **do**
3:     $a \leftarrow \text{next}(i)$
4:     $b \leftarrow \text{next}(j)$
5:     **if** $a = \infty$ and $b = \infty$ **then**
6:         **return** 0
7:     **end if**
8:     **if** $a \neq b$ **then**
9:         **if** $p \oplus [a < b]$ **then**
10:             **return** 1
11:         **else**
12:             **return** -1
13:         **end if**
14:     **end if**
15: **end while**

---

number of symbols in the un-partitioned models, we try to build partitions not exceeding it; when this inevitably happens, a new partition is created.

If the result exceeds the limit of partitions set at the beginning, a new round is performed doubling the quantity of the max symbols. This method repeats until the partition limit is met. The process is fast as the number of allowed symbols doubles at each iteration, hence the number of partitions is typically cut in half or more. The result is partitions that have a limited number of symbols and have similar frequency maps cardinality. Algorithm 4 illustrates the procedure.

Algorithm 4 Partitioning procedure after Gray code ordering.

1: $maxSymbols \leftarrow avergage \ \# \ symbols \ per \ model$
2: $limit \leftarrow n * partitionsPercentage$
3: $splits \leftarrow n$
4: $assignments \leftarrow [-1, \ldots, -1]$
5: **while** $splits > limit$ **do**
6:     $i \leftarrow 0$
7:     $symCounter \leftarrow \emptyset$
8:     **for** $model \in models$ **do**
9:         **if** $|symCounter| > maxSymbols$ **then**
10:             $symCounter \leftarrow \emptyset$
11:             $i \leftarrow i + 1$
12:         **end if**
13:         $symCounter \leftarrow symCounter \cup model.symbols$
14:         $assignments[model] = i$
15:     **end for**
16:     $splits \leftarrow i + 1$
17:     $maxSymbols \leftarrow maxSymbols * 2$
18: **end while**
19: **return** $assignments$

4.3.2 Heuristic escape

After the partitioning procedure, each split is condensed into a single model, merging the frequency maps as introduced at the beginning of the section. The merging procedure should create models that are related to each other, as they are ordered

via Gray-code, but it is safe to assume that the cardinality of those frequency maps is way bigger than a single optimal model for a node. To mitigate this behavior, escaping can be applied to the merged models, this time in a more sophisticated way.

Escaping the aggregated models   To reduce the number of symbols in the merged models we apply escaping before initializing the ANS structures, by minimizing a function. All the frequencies are now considered at the cluster level. The quantity to minimize is:

$$E_{ans} + E_{esc} + escape\ bits + model\ bits$$

The first component is the entropy of the ANS encoding:

$$E_{ans} = -\sum_{s \in S} c_s \log p_s$$

where S is the alphabet, c_s is the count of a symbol, p_s is defined as the frequency of the symbol over the total count of frequencies. This part of the formula models the number of bits required by ANS states.

The second part is the entropy of the escapes, let S' be the set of escaped symbols, $p_{esc} = \sum_{s \in S'} p_s$ is the sum of the probabilities, $c_{esc} = \sum_{s \in S'} c_s$ the sum of counts:

$$E_{esc} = -c_{esc} \log p_{esc}$$

This factor accounts for the bits required to encode the escape symbol into the ANS state. The sum of this and the previous term results in the total bits spent by the ANS encoding.

The third element is the cost of writing escapes on disk, which can be approximated with:

$$escape\ bits = \alpha(c_{esc} * b_{max})$$

where b_{max} is the binary magnitude of the maximum escape, α is the percentage of space gained by using the PFOR-like approach instead of b_{max} bits for each

element. It can be measured empirically, a good value can be 0.9.

The final term of the formula is the space required for the ANS models, that is all the symbols written in γ code, plus the sorted frequencies by gap, this term is referred to as *model bits*.

An important note is that by minimizing this formula, we are considering the symbol statistics at the cluster level, which means, the aggregated symbol counts coming from the nodes. To make a more accurate estimation, one would need to fix a subset S' of escaped symbols, iterate over all the points in the cluster and compute the space required by each node. This would cost too much efficiency-wise.

Choosing the escape symbols set   The escape symbols set S' is the core of the space formula to minimize. The approach does not consider all possible subsets of S, indeed, given the symbols sorted by descending frequency, we fix a cut point and take all the elements after that point as the escape set.

For instance, considering the symbols ordered by frequency in descending order:

$$syms = 1, 2, 5, 3, 13, 110, 45$$

The cut point starts at the end of the symbols, so $S' = \emptyset$, then with each iteration, the last symbol of the list is added, so the set progressively becomes $\{45\}$, $\{45, 110\}$, $\{45, 110, 13\}$ and so on.

This permits minimizing the formula in $O(|S|)$, as all the components can be updated in constant time from one iteration to the other. Figure 4.1 on the following page shows a comparison between the heuristic minimum with respect to the real occupied space.

Escaping before partitioning   The idea of escaping a single node model introduced in the previous section can be used before partitioning. Indeed, the escape wildcard is treated like any symbol in the alphabet during the algorithm. About the sorting criteria when choosing what to escape in the single nodes, preferring smaller symbols is more useful for the clustering procedure, as they are more likely to appear in a wider range of nodes. Lastly, escaping before clustering does not
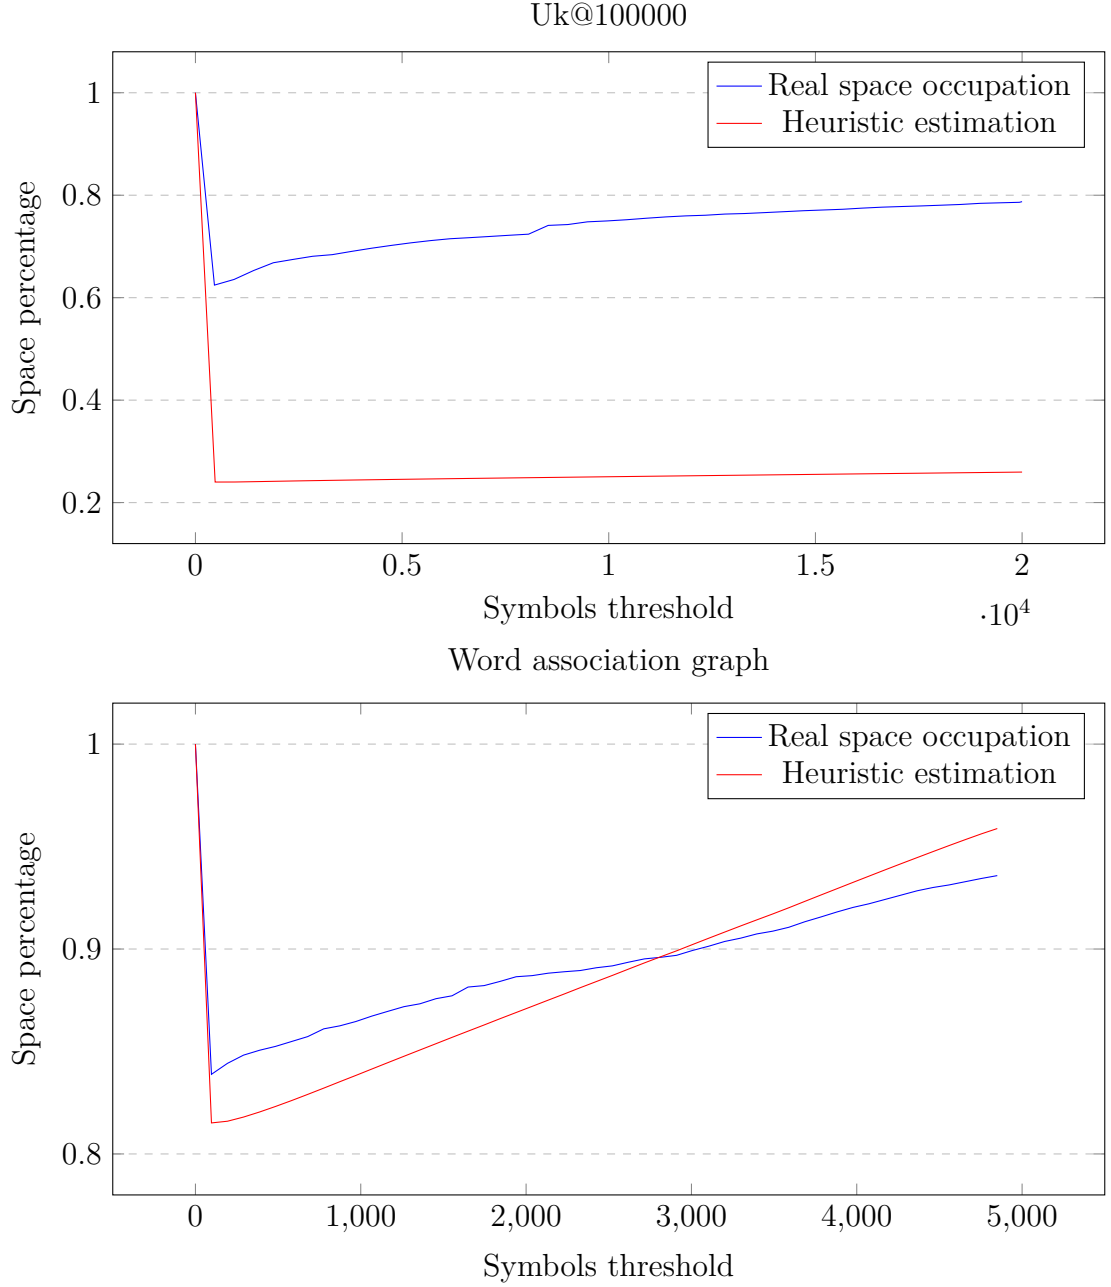
Figure 4.1: The x axis represent the threshold for the escape symbols set, for instance $x = 100$ means that the first 100 symbols are considered in the ANS model, while the remaining $|S| - 100$ are escaped. The blue line is the real space occupied by the graph with that escape set, and the red one show the heuristic function. Both function values are normalized, to have a better visualization. The two minimums for both graphs are close.

lead to the same results as clustering alone, sometimes it worsens the performance, speeding up the method.

Performance This approach limits the excessive size of the *.model* file, while increasing the space taken by the *.graph* stream. On average, clustering saves 40 percent of space compared to the initial approach.

Graph	BV	EF	Naïve ANS	Clustered ANS
cnr-2000	2.897	17.838	17.143	9.152
uk-2007-05@100000	2.033	13.7	9.137	4.736
uk-2007-05@1000000	1.68	16.479	7.849	3.939
en-wiki-2013	13.114	19.500	28.838	16.514
in-2004	2.172	18.981	14.799	7.593
indochina-2004	1.457	18.469	8.572	4.371

Table 4.5: Bits per link obtained by clustered ANS with respect to previous results. The improvement is considerable, clustering saves from 40 to 50 percent of the space occupied by the naïve methodology.

4.4 Implementation details

To have a clearer idea about the implementation of the proposed methodology, we here present the details about the procedures followed when storing and loading a graph on and from disk. It follows an explanation of how successors of a node are retrieved when a request to them is made. The proposed methodology implementation is designed to work seamlessly alongside Webgraph[1], to have an easy time loading and storing graphs. The fastutil[2] and dsiutils[3] Java packages are also heavily used, to have more space and time efficient data structures, and some useful utility functions. All the implementation is publicly available on Github[4].

---

[1] https://github.com/vigna/webgraph
[2] https://github.com/vigna/fastutil
[3] https://github.com/vigna/dsiutils
[4] https://github.com/tomfran/ANS-Graph-compression

4.4.1   Storing and loading graphs

By developing the code in extension of the Webgraph framework, the publicly available compressed graphs can be used to make experiments. In particular, by loading an existing one, we can visit it and store it again with the proposed methodology to assess its goodness.

Re-encoding an existing graph   To store a graph with ANS, we need access to every node successors list. To do so, we can load an existing compressed representation with BVGraph, which is the class defined by the Webgraph framework. To apply for instance the naïve ANS approach one would need to iterate over all the nodes, get the successors, compute statistics and encode them via ANS. Clustering is a different story, as one needs the observed frequency maps in memory to order and partition them.

When the required computation is done, the successors and the models are stored on disk on two separate files, plus an additional one that contains graph properties. This file includes useful information to rebuild the graph, such as the endianness of the integer coding, the number of ANS models, the number of nodes, arcs, and so on. This is also where statistics are stored, such as bits per link, the maximum number of states, and many more. Table 4.6 on the next page shows the properties file for indochina-2004 graph.

Loading from disk   Once a graph is encoded on disk, it can be loaded and visited by decoding the successors of nodes. Given a graph, the property file is used to determine the endianness of the encoding and the number of total models m. Once this is known, m models are allocated in memory loading them from the .model file.

4.4.2   Successors retrieval

The most useful thing to do with a loaded graph is to request the successors of a node. This is straightforward and makes use of the pre-loaded models. Recall that in the .graph file, each node contains the model id, outdegree, and list of

successor states followed by the escaped ones, thus to load node successors a copy of the corresponding pre-loaded model is obtained, then an ANS decoder is built loading all the states and escapes to memory. Successors are then retrieved by performing decode primitives on the states, retrieving escaped symbols when the special character is encountered.

Property	Value
method	cluster
byte order	little endian
nodes	7414866
arcs	194109311
number of models	556
max number of symbols	1503
avg number of symbols	469.182
max number of states	505
avg number of states	1.138
written bits	848512704
bits for states	646330996
bits for escapes	146741003
bits for outdegrees	48464986
bits for successors	793071999
bits for models	6975666
bits per link	4.371
escaped edges	5742212
max number of escapes	1300
escaped edges percentage	0.03
avg number of escapes	0.774

Table 4.6: Properties file for indochina-2004 graph.

Chapter 5

Experimental results

The experimental results are now presented, focusing on the compression ratio and the speed of successors retrieval. The datasets used in the experiments come from the Webgraph archive and are presented in Table 5.1 on the facing page. We find a mixture of social and web graphs of different dimensions. All the graphs are considered in both natural and transposed version.

5.1 Compression

The compression results on all tested graphs can be seen in the tables and plots on the next pages. We can see how the proposed methodology always falls between the succinct representation and the WebGraph framework. Overall it saves 52.07 percent of space compared to the non-compressed representation. The best results are obtained on web graphs, as the average space taken by ANS is 29.19 percent of the succinct bound. Social graphs are a different story, ANS uses on average 74.96 percent compared to Elias-Fano lists.

    The main differences between storing these types of graphs are in the number of escaped edges. Indeed, web graphs present high consecutivity, this means that higher gaps are not frequent in the node successors' gap encoding, leading to lower escaping on the ANS models. Social graphs show less consecutivity, thus higher gaps are more frequent. This means that the heuristic escaping procedure prefers

to remove more symbols, as the space taken by the models would otherwise be too much. More escaping means more overhead on the .graph file, hence performance on social graphs is worse. A detailed analysis of the space occupation can be found in the next paragraph.

Graph	Nodes	Arcs
amazon-2008	735 323	5 158 388
arabic-2005	22 744 080	639 999 458
cnr-2000	325 557	3 216 152
dblp-2011	986 324	6 707 236
enwiki-2013	4 206 785	101 355 853
enwiki-2020	6 047 510	142 691 609
enwiki-2021	6 261 502	150 124 927
enwiki-2022	6 492 490	159 047 205
eu-2005	862 664	19 235 140
eu-2015-host	11 264 052	386 915 963
hollywood-2011	2 180 759	228 985 632
imdb-2021	2 996 317	10 739 291
in-2004	1 382 908	16 917 053
indochina-2004	7 414 866	194 109 311
uk-2002	18 520 486	298 113 762
uk-2007-05@100000	100 000	3 050 615
uk-2007-05@1000000	1 000 000	41 247 159

Table 5.1: Graphs used in the experiments.

Graph	BV		ANS		EF
amazon-2008	9.251	(42.85%)	17.735	(82.14%)	21.591
arabic-2005	1.841	(8.68%)	5.149	(24.29%)	21.201
cnr-2000	2.897	(16.24%)	9.152	(51.31%)	17.838
dblp-2011	8.716	(40.78%)	20.541	(96.10%)	21.374
enwiki-2013	13.114	(67.25%)	16.514	(84.69%)	19.5
enwiki-2020	13.506	(67.83%)	16.872	(84.74%)	19.911
enwiki-2021	13.504	(67.85%)	16.809	(84.45%)	19.904
enwiki-2022	13.518	(67.85%)	16.782	(84.23%)	19.923
eu-2005	3.726	(21.29%)	10.041	(57.36%)	17.504
eu-2015-host	3.273	(20.80%)	7.066	(44.91%)	15.734
hollywood-2011	5.113	(34.92%)	7.719	(52.72%)	14.642
imdb-2021	8.763	(34.18%)	19.6	(76.45%)	25.639
in-2004	2.172	(11.44%)	7.593	(40.00%)	18.981
indochina-2004	1.457	(7.89%)	4.371	(23.67%)	18.469
uk-2002	2.271	(9.98%)	7.314	(32.14%)	22.757
uk-2007-05@100000	2.033	(14.84%)	4.736	(34.57%)	13.7
uk-2007-05@1000000	1.68	(10.19%)	3.939	(23.90%)	16.479

Table 5.2: Bits per link obtained on the natural version of the used graphs. The quantity in parenthesis indicates the space occupied with respect to the quasi-succinct bound set by EF.

Graph	BV		ANS		EF
arabic-2005	1.251	(7.05%)	5.433	(30.62%)	17.744
cnr-2000	2.343	(16.15%)	10.808	(74.52%)	14.504
enwiki-2013	10.954	(64.26%)	14.356	(84.22%)	17.046
enwiki-2020	11.635	(65.44%)	15.112	(84.99%)	17.781
enwiki-2021	11.697	(65.76%)	15.135	(85.09%)	17.788
enwiki-2022	11.684	(65.51%)	15.069	(84.49%)	17.836
eu-2005	2.858	(19.50%)	7.521	(51.32%)	14.656
eu-2015-host	3.578	(22.87%)	7.764	(49.63%)	15.644
in-2004	1.813	(10.83%)	8.677	(51.83%)	16.741
indochina-2004	1.063	(6.69%)	4.923	(30.98%)	15.893
uk-2002	1.703	(8.50%)	8.31	(41.46%)	20.045
uk-2007-05@100000	1.431	(14.49%)	4.51	(45.67%)	9.876
uk-2007-05@1000000	1.046	(9.20%)	3.449	(30.34%)	11.369

Table 5.3: Bits per link obtained on the transposed version of the used graphs. Note that undirected graphs are removed from this table, such as amazon-2008, imdb-2021 and hollywood-2011.
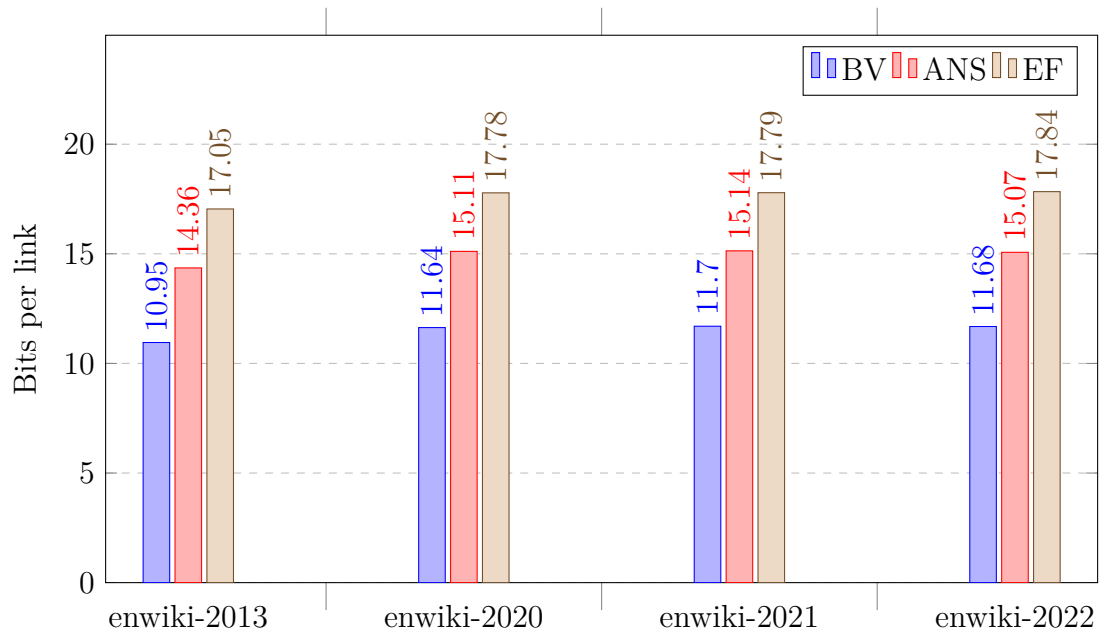
Figure 5.1: Compression ratio for normal wikipedia graphs.
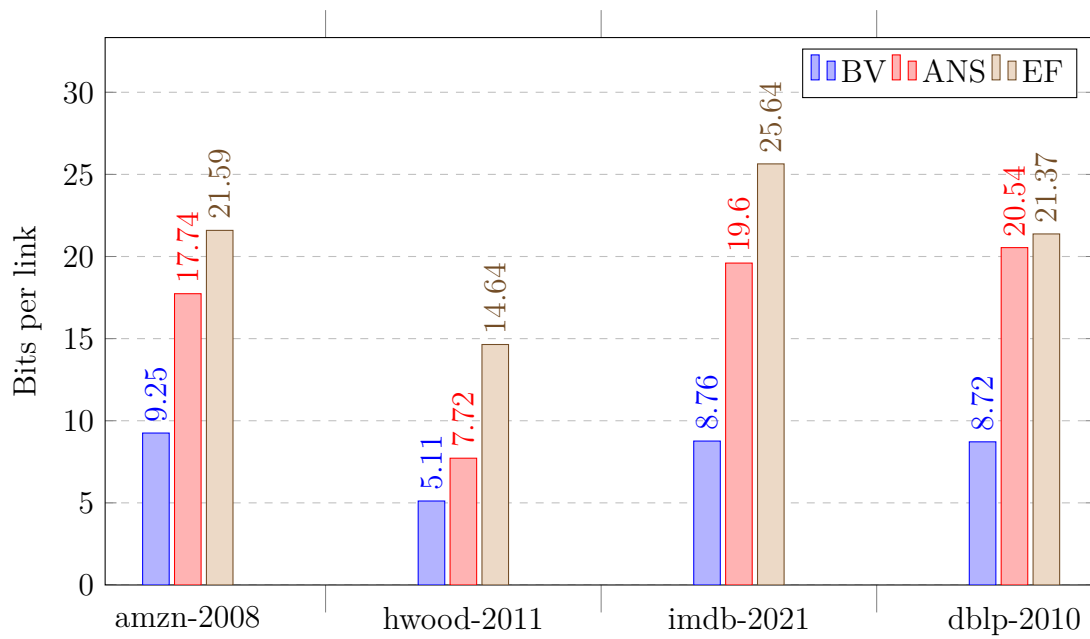

Figure 5.2: Compression ratio for other social graphs

Figure 5.3: Compression ratio for normal smaller web graphs

Figure 5.4: Compression ratio for normal larger web graphs

Space analysis The tables and plots show how compression on web graphs is significantly better than on social ones. By looking into the bits taken by each part of the encoding we can better understand what is happening under the hood when storing these two types of graphs.
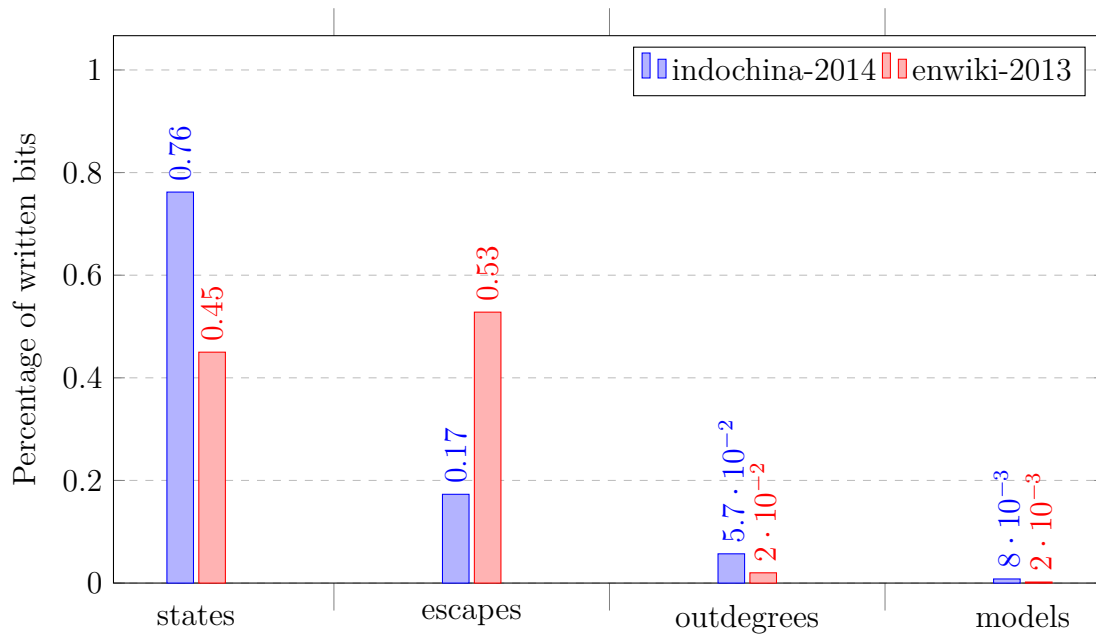

Figure 5.5: Percentage of the total space taken by the models, states and escapes and outdegrees encoding on a web and social graph.

The encoding of successors, namely states plus escapes and outdegrees, completely dominate the space taken by the models for both types of graphs. The plot shows just how good the impact of clustering and escaping is on the .model file. In fact, the models space in the naïve approach was roughly double the successors encoding.

The main difference between the two graphs is the space taken by the escaped edges. In indochina-2014 it accounts for only 17 percent of the total, on the other hand, in enwiki-2013 it accounts for more than half of the total space. This is motivated by the higher escaping percentage obtained on the social graph.

5.2   Speed

To evaluate the speed of the proposed methodology a test is performed for random accesses to node successors. We fix several warmup and measurement rounds, the first ones are to prepare the JVM, and the last ones are where the measurements are made. At each step, a random node is extracted and its successors are visited, this permits finding the time it takes to visit a single node and a single link.

The tests are performed on four graphs, that are uk-2007-05@1000000, indochina-2004, enwiki-2013, and dblp-2011 respectively. The number of warmup iterations is set to two thousand, while five thousand iterations are performed as measurement ones. Each iteration randomly draws fifteen hundred nodes. The measured speed is in terms of nodes and arcs per second, and nanosecond per node and link.

The results depend on how much escaping is performed on the proposed methodology: visiting the web graphs is quite fast, they indeed have little to no escape, at only three percent of the total arcs. Visiting social graphs is slower, the escape percentage for those two graphs is more or less 48, therefore many arcs need to be handled twice, firstly for the state decoding and secondly to read them from the escape symbols list. Overall performance is promising as it is close to BV, EF is way faster than both methods, but no compression takes place there.

Method	nodes/s	arcs/s	ns/node	ns/link
BV	640,837.103	25,966,719.428	1,560.459	38.511
ANS	629,016.732	25,487,757.983	1,589.783	39.235
EF	3,984,096.617	161,435,594.910	250.998	6.194

Table 5.4: Speed results on uk-2007-05@1000000.

Method	nodes/s	arcs/s	ns/node	ns/link
BV	996,213.627	28,147,019.824	1,003.801	35.528
ANS	734,869.732	20,763,009.421	1,360.785	48.163
EF	4,804,775.375	135,754,123.439	208.126	7.366

Table 5.5: Speed results on indochina-2004.

Method	nodes/s	arcs/s	ns/node	ns/link
BV	717,109.048	15,948,505.237	1,394.488	62.702
ANS	523,702.365	11,647,140.603	1,909.482	85.858
EF	5,037,490.309	112,033,784.478	198.512	8.926

Table 5.6: Speed results on enwiki-2013.

Method	nodes/s	arcs/s	ns/node	ns/link
BV	1,735,083.190	11,643,564.924	576.341	85.884
ANS	2,234,998.349	14,998,328.920	447.428	66.674
EF	9,375,914.374	62,918,636.061	106.656	15.894

Table 5.7: Speed results on dblp-2011.

Figure 5.6: Nodes/s over the four tested graphs, the higher the better.

Figure 5.7: Arcs/s on the four tested graphs, the higher the better.

Figure 5.8: Nanoseconds per node on the tested graphs, the lower the better.

Figure 5.9: Nanoseconds per link, the lower the better.

Chapter 6

Conclusions

We have applied asymmetric numeral systems to graph compression, improving performance through various iterations. The first approach is quite rough but, after a lot of refinements and optimizations, the final compression scheme turned out to be competitive. We developed two heuristic approaches to merge a large quantity of ANS models and to add escaping to them, minimizing a space estimation function.

The proposed methodology creates a new tradeoff between space and speed concerning quasi-succinct representations. Indeed, at the cost of slower access speed, the clustered ANS compression scheme can save as much as 76 percent of the space taken by Elias-Fano lists.

The compression ratio obtained by BV graph is superior, but keep in mind that reference encoding is not exploited in the proposed methodology. Another point is that the proposed compression format is agnostic to gap distribution, whereas classical techniques like instantaneous codes rely on the fact that the gap distribution is monotonically decreasing.
The speed of access between the two compression schemes is comparable.

Future works Numerous techniques could be employed to improve compression and the first one is reference encoding. The idea would be to use the same technique as BV graph, adding ANS encoding only for the residuals. This could potentially

improve the first method, as consecutivity in the residuals would create a high probability for smaller gaps.

The focus of the presented methodology is mainly on the reduction of the space taken by the models, perhaps a different approach could be employed. The first idea that comes to mind is using different ordering criteria for the models, and potentially another partition or clustering technique. Another improvement could be exploiting reference encoding for the models on disk, as Gray-code ordered models should share a lot of common symbols. Finally, a preprocessing on the integers to encode can reduce the alphabet size, for instance, applying variable byte encoding to the gaps would create k bits blocks, limiting the alphabet to 2^k symbols. This would probably create smaller representations for the models, wasting space for the successors, as an integer requires more than one encoding.

Bibliography

[1]   Paolo Boldi, Massimo Santini, and Sebastiano Vigna. "Permuting Web and Social Graphs". In: *Internet Mathematics* 6 (Jan. 2009), pp. 257–283. DOI: 10.1080/15427951.2009.10390641.

[2]   Paolo Boldi and Sebastiano Vigna. "The WebGraph framework I: compression techniques". In: *WWW '04*. 2004.

[3]   Paolo Boldi and Sebastiano Vigna. "The WebGraph framework II: codes for the World-Wide Web". In: *Data Compression Conference, 2004. Proceedings. DCC 2004*. 2004, pp. 528–. DOI: 10.1109/DCC.2004.1281504.

[4]   Paolo Boldi et al. "BUbiNG: Massive Crawling for the Masses". In: *Proceedings of the Companion Publication of the 23rd International Conference on World Wide Web*. International World Wide Web Conferences Steering Committee, 2014, pp. 227–228.

[5]   Paolo Boldi et al. "Layered Label Propagation: A MultiResolution Coordinate-Free Ordering for Compressing Social Networks". In: *Proceedings of the 20th international conference on World Wide Web*. Ed. by Sadagopan Srinivasan et al. ACM Press, 2011, pp. 587–596.

[6]   Jarek Duda. "Asymmetric numeral systems". In: *CoRR* abs/0902.0271 (2009). arXiv: 0902.0271. URL: http://arxiv.org/abs/0902.0271.

[7]   Jarek Duda. "Asymmetric numeral systems as close to capacity low state entropy coders". In: *CoRR* abs/1311.2540 (2013). arXiv: 1311.2540. URL: http://arxiv.org/abs/1311.2540.

[8]   Jarek Duda et al. "The use of asymmetric numeral systems as an accurate replacement for Huffman coding". In: *2015 Picture Coding Symposium (PCS)* (2015), pp. 65–69.

[9]   Peter Elias. "Efficient Storage and Retrieval by Content and Address of Static Files". In: *J. ACM* 21.2 (Apr. 1974), pp. 246–260. ISSN: 0004-5411. DOI: `10.1145/321812.321820`. URL: `https://doi.org/10.1145/321812.321820`.

[10]   Peter Elias. "Universal codeword sets and representations of the integers". In: *IEEE Trans. Inf. Theory* 21 (1975), pp. 194–203.

[11]   Solomon W. Golomb. "Run-length encodings (Corresp.)" In: *IEEE Trans. Inf. Theory* 12 (1966), pp. 399–401.

[12]   Keith Henderson, Brian Gallagher, and Tina Eliassi-Rad. "EP-MEANS: An Efficient Nonparametric Clustering of Empirical Probability Distributions". In: *Proceedings of the 30th Annual ACM Symposium on Applied Computing.* SAC '15. Salamanca, Spain: Association for Computing Machinery, 2015, pp. 893–900. ISBN: 9781450331968. DOI: `10.1145/2695664.2695860`. URL: `https://doi.org/10.1145/2695664.2695860`.

[13]   D. Lemire and L. Boytsov. "Decoding billions of integers per second through vectorization". In: *Software: Practice and Experience* 45.1 (May 2013), pp. 1–29. DOI: `10.1002/spe.2203`. URL: `https://doi.org/10.1002%2Fspe.2203`.

[14]   Antonio Mallia, Michal Siedlaczek, and Torsten Suel. "An Experimental Study of Index Compression and DAAT Query Processing Methods". In: *Advances in Information Retrieval - 41st European Conference on IR Research, ECIR 2019, Cologne, Germany, April 14-18, 2019, Proceedings, Part I.* Ed. by Leif Azzopardi et al. Vol. 11437. Lecture Notes in Computer Science. Springer, 2019, pp. 353–368. DOI: `10.1007/978-3-030-15712-8\_23`. URL: `https://doi.org/10.1007/978-3-030-15712-8%5C_23`.

[15]   Alistair Moffat and Matthias Petri. "ANS-Based Index Compression". In: *Proceedings of the 2017 ACM on Conference on Information and Knowledge Management.* CIKM '17. Singapore, Singapore: Association for Com-

puting Machinery, 2017, pp. 677–686. ISBN: 9781450349185. DOI: `10.1145/3132847.3132888`. URL: `https://doi.org/10.1145/3132847.3132888`.

[16] Alistair Moffat and Matthias Petri. "Index Compression Using Byte-Aligned ANS Coding and Two-Dimensional Contexts". In: *Proceedings of the Eleventh ACM International Conference on Web Search and Data Mining*. WSDM '18. Marina Del Rey, CA, USA: Association for Computing Machinery, 2018, pp. 405–413. ISBN: 9781450355810. DOI: `10.1145/3159652.3159663`. URL: `https://doi.org/10.1145/3159652.3159663`.

[17] Giulio Ermanno Pibiri and Rossano Venturini. "Techniques for Inverted Index Compression". In: *ACM Computing Surveys* 53.6 (Nov. 2021), pp. 1–36. DOI: `10.1145/3415148`. URL: `https://doi.org/10.1145%2F3415148`.

[18] L.H. Thiel and H.S. Heaps. "Program design for retrospective searches on large data bases". In: *Information Storage and Retrieval* 8.1 (1972), pp. 1–20. ISSN: 0020-0271. DOI: `https://doi.org/10.1016/0020-0271(72)90024-1`. URL: `https://www.sciencedirect.com/science/article/pii/0020027172900241`.

[19] Tai Vo Van and T. Pham-Gia. "Clustering probability distributions". In: *Journal of Applied Statistics* 37.11 (2010), pp. 1891–1910. DOI: `10.1080/02664760903186049`. eprint: `https://doi.org/10.1080/02664760903186049`. URL: `https://doi.org/10.1080/02664760903186049`.

[20] Sebastiano Vigna. "Quasi-Succinct Indices". In: (2012). DOI: `10.48550/ARXIV.1206.4300`. URL: `https://arxiv.org/abs/1206.4300`.

[21] Hao Yan, Shuai Ding, and Torsten Suel. "Inverted index compression and query processing with optimized document ordering". In: Jan. 2009, pp. 401–410. DOI: `10.1145/1526709.1526764`.